HUMBOLDT UNIVERSITÄT ZU BERLIN

SEMINAR PAPER

Numerical Methods for solving Eigenvalue-Problems

Thomas Siskos (580726)

NUMERICAL INTRODUCTORY COURSE

Supervised by:

Prof. Dr. Brenda López Cabrera

July 1, 2018

Contents

List of Algorithms

List of Figures

List of Tables

1   Motivation

Abstract

Eigenvalues and eigenvectors are often the solution to multidimensional optimization problems, however computing them by hand for anything but trivial matrices is most of the time infeasible or inpractical. To this extend we would like to deploy an automated procedure which yields the correct eigenvectors and eigenvalues. We demonstrate the relevance of eigenvalues and eigenvectors by revising two applications from statistics, Principal Component Analysis and Fisher's Linear Discriminant Analysis, which we follow up by investigating four algorithms suited for eigenvalue problems. Finally we provide a compound solution that takes advantage of each algorithms strengths.

For many statistical applications eigenvectors provide a formidable solution. Be it dimensionality reduction in terms of a Principal Component Analysis or classification by Fisher's Linear Discriminant Analysis, both come in the guise of optimization problems. But what are eigenvalues and eigenvectors?

If A is an $n \times n$ matrix, v is a non-zero vector and λ is a scalar, such that

$$Av = \lambda v \tag{1}$$

then v is called an *eigenvector* and λ is called an *eigenvalue* of the matrix A. An eigenvalue of A is a root of the characteristic equation,

$$det\left(A - \lambda I\right) = 0. \tag{2}$$

Geometrically speaking, we require a vector which, when multiplied by matrix A, will not get rotated but only elongated by a factor λ.

When confronted with a high-dimensional data matrix $X \in \mathbb{R}^{n \times m}$ an analyst often wishes to find a lower-dimensional representation, while conserving as much of the structure as possible. One way of achieving this goal is to choose a standardized linear combination of features that aim to maximize the variance of the projection $\delta'X$. We can formalize this as

$$max \ \delta'Var\left(X\right)\delta \ s.t. \ \sum \delta_i^2 = 1. \tag{3}$$

The solution to this optimization problem is given by

$$Y = \Gamma' (X - \mu) \tag{4}$$

where $Y \in \mathbb{R}^{n \times m}$ is the matrix of rotations, $\Gamma \in \mathbb{R}^{m \times m}$ is the matrix of eigenvectors, $\mu \in \mathbb{R}^m$ is the vector of sample means. [Härdle and Simar, 2015]

A second use for Eigenvalues lies in Fisher's Linear Discriminant Analysis for classification problems. In this setting we would like to find the rotation that is best suited to separate the classes. To this effect we try to maximize the between class variance while simoultaneously trying to minimize the within-class variance. This is equivalent to maximizing the so called Rayleigh-coefficient

$$max J = \frac{J_b(w)}{J_w(w)} = \frac{w\prime S_b w}{w\prime S_w w}. \tag{5}$$

The matrices for between and within class variance are defined as $S_b = \sum_{c=1}^{C}(\mu_c - \mu)(\mu_c - \mu)\prime$ and $S_w = \sum_{c=1}^{C} \sum_{i \in c}(x_i - \mu_c)(x_i - \mu_c)\prime$

The solution to this maximization problem is given by

$$S_B S_W S_B v = \lambda v, \tag{6}$$

which again is an eigenvalue problem [Duda et al., 2000].

In the following we will try to extract the eigenvalues of matrices by reformulating any complicated eigenvalue problem into a straightforward one by diagonalizing the matrix in question, without altering the eigenvalues we would like to compute. In section two we provide the mathematical tools, which allow us to do so. We follow these justifications by proposing two main algorithms for computing eigenvalues, first the Jacobi-Method for symmetric matrices, then the QR-Method for arbitrary square matrices in section 3. Additionally, for the QR-Method we define two extensions which try to increase the initial QR-algorithm's speed. For all algorithms we provide implementations in the `Python`-programming-language, all plots have been produced with the `matplotlib`-module and the custom algorithms rely solely on modules from `Python`'s `scipy` environment for sci-

entific computing [van Rossum, 1995, Hunter, 2007, McKinney, 2010]. In section 4 we will analyse the implemented routines by critically reflecting upon the accuracy of the obtained results as well as their efficiency. In the final section we provide a final algorithm which combines the strengths of the defined procedures by chosing the algorithm that is most fit for the underlying problem.

2 Similarity Transformations

In general we want to reformulate the eigenvalue problem of a complicated matrix into an eigenvalue problem of a simple matrix, which yields the same eigenvalues. Simple matrices in our case will be diagonal matrices, since with them it is possible to identify their eigenvalues simply as entries on the main diagonal. Such a transformation that conserves the eigenvalues of a matrix is called a *similarity transformation.*

Two $n \times n$ matrices A and B are called *similar* if there exists an invertible matrix P such that

$$A = P^{-1}BP. \tag{7}$$

It is obvious that the similarity relationship is commutative as well as transitive. If A and B are similar, it holds that

$$B - \lambda I = P^{-1}BP - \lambda P^{-1}IP$$

$$= A - \lambda I.$$

Hence A and B have the same eigenvalues. This fact also follows immediately from the transitivity of the similarity relationship and the fact that a matrix is similar to the diagonal matrix formed from its eigenvalues, as stated in the spectral-decomposition. Important types of similarity transformations are based around orthogonal matrices. If Q is orthogonal and

$$A = Q'BQ,$$

A and B are called *orthogonally similar* [Gentle, 1998]. We will use *orthogonal similarity transformations* to diagonalize matrices we wish to know the eigenvalues of.

3 Algorithms

3.1 Jacobi Method

The *Jacobi-Method* for computing the eigenvalues of a symmetric matrix $A \in \mathbb{R}^{n \times n}$ deploys a sequence of orthogonal similarity transformations that eventually results in

$$A = P\Lambda P^{-1} \Leftrightarrow \Lambda = P^{-1}AP,$$

where Λ is diagonal and P consists of a sequence of matrix multiplications $P = \prod_{k=1}^{K} V_{p_k,q_k}(\theta_k)$ and $V_{p_k,q_k}(\theta_k)$ is of the form proposed in (20). More specific the *Jacobi iteration* is

$$A^{(k)} = V'_{p_k,q_k}(\theta_k)A^{(k-1)}V_{p_k,q_k}(\theta_k), \tag{8}$$

where p_k, q_k and θ_k are chosen such that A^k resembles more a diagonal matrix than A^{k-1}. Specifically they will be chosen as to reduce the sum of squares of the off-diagonal elements. As we saw in (20) it is easy to chose an angle θ_k in order to introduce a zero in a single Givens rotation. Here we use the rotations in the context of a similarity transformation, so it is a little more complicated.

We require that $a_{pq}^{(k)} = 0$, this implies

$$a_{pq}^{(k-1)}(\cos^2 \theta - \sin^2 \theta) + \left(a_{pp}^{(k-1)} - a_{qq}^{(k-1)}\right)\cos \theta \sin \theta = 0. \tag{9}$$

Figure 1: Progress Jacobi-Method

We can use the trigonometric identities

$$\cos(2\theta) = \cos^2\theta \sin^2\theta$$

$$\sin(2\theta) = 2\cos\theta\sin\theta,$$

in (9) we have

$$\tan(2\theta) = \frac{2a_{pq}^{(k-1)}}{a_{pp}^{(k-1)} - a_{qq}^{(k-1)}}.$$

From this we can retrieve the angle and obtain the rotation matrix in each iteration [Gentle, 1998].

The algorithm converges if the off-diagonal elements are sufficiently small. The best index pair at a given iteration is the pair (p, q) that satisfies

$$|a_{pq}^{(k-1)}| = \max_{i<j}|a_{ij}^{(k-1)}|.$$

If this choice is made, the Jacobi Method can be shown to converge [Gentle, 1998].

Figure 1 visualizes the progress of the *Jacobi*-method on a symmetric 5×5 matrix. As we can see, in the first iteration the element $a_4 3$ is eliminated. In the subsequent operations the *Jacobi*-method continues to eliminate any non-zero entries on the off-diagonal until the algorithm convergences after 10 iterations.

Algorithm 1 `jacobi`

Require: symmetric matrix A

Ensure: $0 < precision < 1$

     initialize: $L \leftarrow A$; $U \leftarrow I$; $L_{max} \leftarrow 1$

1: **while** $L_{max} > precision$ **do**

2:       Find indices i, j of largest value in lower triangle of $abs(L)$

3:       $L_{max} \leftarrow L_{i,j}$

4:       $\alpha \leftarrow \frac{1}{2} \cdot \arctan(\frac{2A_{i,j}}{A_{i,i} - A_{j,j}})$

5:       $V \leftarrow I$

6:       $V_{i,i}, V_{j,j} \leftarrow \cos\alpha$; $V_{i,j}, V_{j,i} \leftarrow -\sin\alpha, \sin\alpha$

7:       $A \leftarrow V'AV$; $U \leftarrow UV$

8: **end while**

9: **return** $diag(A)$, U

3.2 QR-Method

The most widely used algorithm to extract eigenvalues is the so called *QR*-method. The most important advantage of the *QR*-method over the *Jacobi*-method is that it can be applied to non-symmetric matrices. Note however, that it is simpler for symmetric matrices, since the eigenvalues are real-valued.
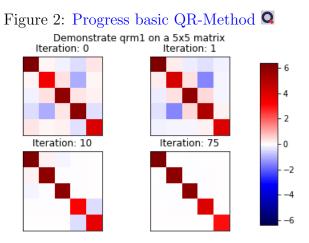
The *QR*-method to extract the eigenvalues of a square matrix $A \in \mathbb{R}^{n \times n}$ is performed by first computing the titular QR decomposition of A.

$$A = QR, \tag{10}$$

where Q is an orthogonal and R is an upper triangular matrix. Then define the QR iteration as

$$A^k = Q'_{k-1} A_{k-1} Q_{k-1} = R_{k-1} Q_{k-1} \tag{11}$$

Note hereby that all matrices in the sequence $\{A_k\}$ share the same eigenvalues, since this procedure is a similarity transformation due to Q's orthogonality. Additionally, in an implementation it is usually preferable to compute the QR-iteration in the way shown at the rightmost part of equation (11) [Börm and Mehl, 2012]. Although, mathematically, each statement is exactly identical there is a practical difference in due to computational imperfections and limited machine precision. The reason being that the computation

Figure 2: Progress basic QR-Method

of $Q'_{k-1}A_{k-1}Q_{k-1}$ obviously requires two matrix multiplications whereas the result of $R_{k-1}Q_{k-1}$ can be readily obtained by one. When combining multiple steps over a long sequence of QR iterations the additional computations lead to additional rounding errors, which can have an influence on the accuracy of the obtained results. Additionally, less computations lead of course to a faster procedure in general.

Figure 3.2 visualizes the progress of the basic QR-method on the same 5×5 matrix as in Figure 1. Compared to the *Jacobi*-method it does not explicitly pick a single element that will be eliminated per iteration. Instead, the QR-method extracts the eigenvalues by a process that is called "chasing". By that we mean that alternating steps are being performed, which create non-zero eintries in positions $(i+2, i)$, $(i+3, i)$ and $(i+3, i+1)$ and restore them to zero, as the nonzero entries are moved farther down the matrix [Gentle, 1998]. We can also see that compared to the *Jacobi*-Method, so far, the QR-algorithm lacks in speed. Where the *Jacobi*-method was almost done diagonalizing the matrix in iteration 10, the basic QR-algorithm still had multiple non-zero entries left. Thus we would like to make minor improvements on the algorithm's efficiency.

3.2.1 Hessenberg Variant

In order to speed up the QR-method it is advisable to transform the matrix to its upper *Hessenberg* form. A matrix A is of upper *Hessenberg* form if it is upper triangular except

---

Algorithm 2 `QRM1` 🔴

Require: square matrix A
    initialize: $conv \leftarrow False$
 1: **while** not $conv$ **do**
 2:    $Q, R \leftarrow$ QR-Factorization of A
 3:    $A \leftarrow RQ$
 4:    **if** A is diagonal **then**
 5:        $conv \leftarrow$ `True`
 6:    **end if**
 7: **end while**
 8: **return** $diag(A), Q$

---

for the first subdiagonal, which may be non-zero. In particular $a_{ij} = 0 \; \forall i > j + 1$:

$$\begin{bmatrix} X & X & X & \ldots & X & X \\ X & X & X & \ldots & X & X \\ 0 & X & X & \ldots & X & X \\ 0 & 0 & X & \ldots & X & X \\ \vdots & \vdots & & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \ldots & X & X \end{bmatrix}$$

A matrix can be reduced to *Hessenberg* form in a finite number of similarity transformations using Householder transformations or Givens rotations. For symmetric matrices the transformation into a *Hessenberg*-form results in a tridiagonal matrix. But even for non-symmetric matrices, the *Hessenberg*-form allows a large saving in subsequent computations. After the transformation we can deploy the previously defined *QR*-method [Gentle, 1998].

Figure 3 visualizes the progress of the *Hessenberg* variant of the *QR*-method. In order to make it comparable to the previous algorithms the same 5×5-dimensional matrix is being evaluated. We can readily see, that the transformation to *Hessenberg*-form results in a tridiagonal matrix. This facilitates computations and explains the vastly improved resulting matrix after 10 iterations compared to the basic *QR*-method. However, it still does not match the progress of the *Jacobi*-method after the same number of iterations.

Algorithm 3 QRM2

Require: square matrix A

1: $A \leftarrow \texttt{hessenberg}(A)$

2: continue with: QRM1(A)

3.2.2 Accelerated Variant

We could already improve the QR-method and cut down on computational cost. However, we still cannot match the results of the *Jacobi* method. To this effect we present the final adjustment on the QR-method to improve convergence speed. The general idea is, that we deliberately create an additional zero entry on the main diagonal by subtracting a scalar on each element, perform the QR-iteration and finally undo the subtraction. In particular, we define

$$
T^m = \begin{bmatrix}
\alpha_1^m & \beta_1^m & 0 & 0 & \dots & & 0 \\
\beta_1^m & \alpha_2^m & \beta_2^m & & & & \\
0 & \beta_2^m & \alpha_3^m & \beta_3^m & & & \vdots \\
& & \ddots & \ddots & \ddots & & \\
& & & & \beta_{n-2}^m & \alpha_{n-1}^m & \beta_{n-1}^m \\
0 & & & & & \beta_{n-1}^m & \alpha_n^m
\end{bmatrix}
\tag{12}
$$

Figure 3: Progress Hessenberg-QR-Method

$$T^m = T - t_{n,n}I$$

$$T^m = QR$$

$$T^{m+1} = T^m + t_{n,n}I$$

After this we can define the accelerated iteration step as

$$R_m = Q'_m \left(T_m - \alpha_n^m I \right) \tag{13}$$

$$T_{m+1} = Q'_m \left(T_m - \alpha_n^m I \right) Q_m + \alpha_n^m I \tag{14}$$

$$= Q'_m T_m Q_m \tag{15}$$

Again T_{m+1} is similar to T_m.

Figure 4 visualizes the progress of the accelerated QR-method. For comparability, the matrix used is the same 5×5 matrix as before. Most notably is, that the accelerated method still performs worse than the *Jacobi* method. So far results seem similar compared to the *Hessenberg* variant of the QR-method. The case can be made that the accelerated method performs considerably better than the basic QR-method and slightly better than the *Hessenberg* variant after 10 iterations than the. Howbeit, this claim warrants further analysis.

Figure 4: Progress Accelerated QR-Method

Algorithm 4 QRM3 Q

Require: square matrix $A \in \mathbb{R}^{p \times p}$
1: $T \leftarrow$ hessenberg(A), $conv \leftarrow False$
2: **while** not $conv$ **do**
3:     $Q, R \leftarrow$ QR-Factorization of $T - t_{p-1,p-1}I$
4:     $T \leftarrow RQ + t_{p-1,p-1}I$
5:     **if** T is diagonal **then**
6:         $conv \leftarrow True$
7:     **end if**
8: **end while**
9: **return** $diag\left(T\right), Q$

---

4   Analysis

In order to test the quality of our implemented algorithms we will follow a two pronged approach. First, we will scrutinize the quality of the obtained results. Once we can be sure that the implemented routines are reliable or, to be more precise, under which conditions our routines are reliable we can analyze their convergence behavior and compare the efficiency accross algorithms in terms of computational cost.

4.1   Accuracy

If we wish to investigate the accuracy of our algorithms, we first need to define what we mean by that. We set up an environment of unit tests that compare the absolute difference between the vector of true eigenvalues of a matrix and the vector of computed eigenvalues by our custom algorithms. If any element of this L_1-norm is larger than a threshold of 10^{-5} we consider the test as failed. In order to obtain matrices with known eigenvalues we will reverse the spectral decomposition, which states that every square, symmetric matrix $A \in \mathbb{R}^{m \times m}$ can be written as

$$A = \Gamma\Lambda\Gamma' = \sum_{j=1}^{m} \lambda_j \gamma_j \gamma_j', \tag{16}$$

where $\Lambda = diag(\lambda_1, \ldots, \lambda_m)$ is a diagonal matrix with the eigenvalues of A on its main diagonal and where $\Gamma = (\gamma_1, \ldots, \gamma_m)$ is the matrix containing the associated eigenvectors [Härdle and Simar, 2015]. This means we can draw a random vector of eigenvalues λ_{true}

Table 1: Failed unit tests accross matrix-sizes

Algorithm	Maximum number of iterations	Dimension				
		3	4	5	6	7
Jacobi	-	0	0	0	0	0
QRM1	10	785	952	994	999	1000
	100	110	224	349	474	565
	1000	7	16	33	38	45
	10000	1	1	1	1	3
	100000	0	0	0	0	0
QRM2	10	826	975	997	1000	1000
	100	111	191	360	467	602
	1000	9	16	22	35	58
	10000	0	0	3	2	0
	100000	0	0	0	0	0
QRM3	10	251	616	896	983	997
	100	29	52	116	222	311
	1000	6	2	8	19	21
	10000	0	0	1	1	3
	100000	0	0	0	0	0

and a random orthogonal matrix to construct a test matrix A_{test}. Once the matrix A_{test} is initialized, we can plug it into the eigenvalue-routine we wish to test.

The number of failures out of 1000 unit tests on matrices of different dimensions are catalogued in Table 1. Most notably, we see that the *Jacobi*-method outperforms any of the *QR*-methods in terms of accuracy. This is mostly due to its convergence criterion being defined differently than the convergence criterion of the *QR*-methods. Simply put, the *Jacobi*-method will always converge in this setup. Also keep in mind, that since the *Jacobi*-method is only allowed for symmetric matrices, it essentially eliminates two non-zero off-diagonal elements at once. Comparing only the *QR*-methods amongst each other one can clearly see, that the quality of the obtained results increases, the more willing an analyst is to wait. Convergence is slower for the *QR*-methods and a premature stop leads to poor eigenvalues. One the other hand, if the *QR*-methods converge, they unanimously pass all tests. It would therefore be possible to always achieve an arbitrary grade of precision by setting the maximum number of iterations to infinity.

Figure 5: Unit-tests: Iterations 🔍

4.2 Efficiency

Now that we can be confident that each of the routines is able to offer accurate results we are interested in which method is most appropriate for which kind of matrix. Especially, we are interested in convergence speed. Figure 5 shows the distribution of steps performed before the algorithms reached convergence or their maximum allowed number of iterations. Again, the *Jacobi*-method outperforms every *QR*-method. For matrices of small dimensions, i.e. eleven or less, the accelerated *QR*-method outperforms the other *QR*-methods. However, for larger matrices the *Hessenberg* variant seems to be slightly more appropriate, since the accelerated method tends to have a wider range of needed steps. Additionally, the accelerated method's distribution has a wider tail whenever the dimension of the test matrices is larger than eleven. This means that high waiting times, operating on the same matrix, are more likely to occur using the *accelerated* variant compared to the *Hessenberg* variant or even the basic variant. Surprisingly, there seems to be no striking difference between the distributions of the basic and the *Hessenberg* variants. Looking closely, the only effect of performing a *Hessenberg* transformation prior to running a *QR*-algorithm on a test matrix is to mitigate outliers in terms of computation steps needed before convergence. And even that is true, only for large matrices. On the other hand, it does not seem to incurr systematically higher costs.
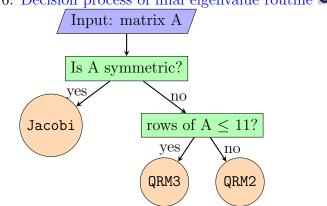
Figure 6: Decision process of final eigenvalue routine

5   Conclusion

We have demonstrated the practical importance of eigenvalues in the context of Principal Component Analysis and Fisher's Linear Discriminant Analysis. For computing the Eigenvalues we have shown various numerical procedures. We examined the *Jacobi*- and the *QR*-method, as well as a selection of simple improvements towards the basic *QR*-method. Finally, we would like to present a compound solution which combines the strengths of all presented algorithms and automatically picks the procedure that fits best to the problem. For obvious reasons, it is best to chose the *Jacobi*-method, whenever applicable. That is, when the input-matrix is symmetric. Whenever the matrix is not symmetric, a second decision based on the matrix' size will determine which *QR*-method will be used. If the matrix' dimension is smaller than eleven it will chose the accelerated *QR* method. Otherwise, the *Hessenberg* variant of the *QR*-method will be deployed.

6   Appendix

6.1   Householder-Reflections

Our goal is still to diagonalize a matrix in order to programmatically extract its eigen-values. So far we have seen that there exist such transformations that conserve the eigenvalues of a given matrix. However, we require transformations that, additionally, eliminate non-zero entries on the off-diagonal elements of said matrix. A greedy tech-nique, that eliminates all but the first elements of a vector is proposed in the form of Householder-Reflections.

Let u and v be orthonormal vectors and let x be a vector in the space spanned by u and v, such that

$$x = c_1 u + c_2 + v$$

for some scalars c_1 and c_2. The vector

$$\tilde{x} = -c_1 u + c_2 v$$

is a *reflection* of x through the line difined by the vector u. Now consider the matrix

$$P = I - 2uu'. \tag{17}$$

Note that

$$Px = c_1 u + c_2 v - 2c_1 uuu' - 2c_2 vuu'$$

$$= c_1 u + c_2 v - 2c_1 u'uu - 2c_2 u'vu$$

$$= -c_1 u + c_2 v$$

$$= \hat{x}.$$

The matrix P is called a reflector. The usefulness of Householder-Reflections stems from the fact that it is easy to transform a vector of the form

$$x = (x_1, x_2, \ldots, x_n)$$

into a vector

$$\hat{x} = (\hat{x}_1, 0, \ldots, 0).$$

If $Qx = \hat{x}$, then $||x||_2 = ||\hat{x}||_2$ and thus $\hat{x}_1 = \pm||x||_2$, since it is the only non-zero entry. To construct the reflector let

$$v = (x_1 + sign(x_1)||x||_2, x2, \ldots, x_n) \tag{18}$$

and $u = \frac{v}{||v||_2}$ [Gentle, 1998]. We use the *sign*-function, which simply returns the sign of its argument in order to avoid the numerical problem known as *catastrophic cancellation.* It can occur when adding two very close, but different, floating point numbers of differing signs. In some unfortunate cases both of these numbers get represented by the same computer number and, because of their opposing signs cancel each other out. In our case this would mean, that we reflect the vector onto the origin. Fortunately, by making use of the *sign* function we can make sure that both summands will share the same sign, thus mitigating any concerns about catastrophic cancellation.

We use reflectors to compute the so called QR factorization of an aribitrary square matrix $A \in \mathbb{R}^{n \times n}$.

$$A = QR \tag{19}$$

where Q is orthogonal and R is upper triangular. We use Householder transformations to reflect the i^{th} column and produce zeros below the (i, i) element. The QR-factorization of a matrix $A \in \mathbb{R}^5$ would therefore consist of five Householder-reflections with $Q = P_5 P_4 P_3 P_2 P_1$. The number of computations for the QR factorization in this fashion is $2n^3/3$ multiplications and $2n^3/3$ additions [Gentle, 1998].

6.2   Givens-Rotations

Another way of forming the QR-factorization is by using orthogonal transformations which rotate a vector in a way such that a specific element becomes 0 and only one other element in the vector being changed. These transformations are called *Givens transformations, Givens Rotations* or *Jacobi transformations*

Figure 7: Rotation of x

A Givens rotation in $\mathbb{R}^2$.

Using orthogonal transformations we can also rotate a vector in such a way that a specified element becomes 0 and only one other element in the vector is changed. The basic idea can be seen in a two-dimensional space. We wish to rotate the vector $x = (x_1, x_2)$ to $\tilde{x} = (\tilde{x_1}, 0)$ as with a reflector.

It is easy to see that the orthogonal matrix

$$Q = \begin{bmatrix} \cos\theta & \sin\theta \\ -\sin\theta & \cos\theta \end{bmatrix}$$

performs the desired rotation, if $\cos\theta = \frac{x_1}{||x||_2}$ and $\sin\theta = \frac{x_2}{||x||_2}$

In general, we can construct an orthogonal $matrix V_{pq}$, that will transform the vector

$$x = (x_1, \ldots, x_p, \ldots x_q, \ldots, x_n)$$

to

$$\tilde{x} = (x_1, \ldots, \tilde{x}_p, \ldots 0, \ldots, x_n)$$

. The matrix that does this is

$$V_{pq}(\theta) = \begin{bmatrix} 1 & & & & & & \\ & \ddots & & & & & \\ & & \cos\theta & & \sin\theta & & \\ & & & \ddots & & & \\ & & -\sin\theta & & \cos\theta & & \\ & & & & & \ddots & \\ & & & & & & 1 \end{bmatrix} \tag{20}$$

where $\cos\theta = \frac{x_p}{||x||}$ and $\sin\theta = \frac{x_q}{||x||}$.

A rotation matrix is therefore the same as an identitiy matrix, in which we change four elements [Börm and Mehl, 2012]. We will use Givens rotations primarily in the Jacobi-Method.

6.3   Eigenvalue Routines

```python
import numpy as np
import copy


def hreflect1D(x):
    """
    Calculate Householder reflection: Q = I - 2*uu'.

    Parameters:
        X: numpy array.

    Returns:
        Qx: reflected vector.
        Q: Reflector (matrix).
    """
    # Construct v:
    v = copy.deepcopy(x)
```

```python
18      v[0] += np.linalg.norm(x)

19

20      # Construct u: normalize v.
21      vnorm = np.linalg.norm(v)
22      if vnorm:
23          u = v / np.linalg.norm(v)
24      else:
25          u = v

26

27      # Construct Q:
28      Q = np.eye(len(x)) - 2 * np.outer(u, u)
29      Qx = np.dot(Q, x)

30

31      return Qx, Q

32

33

34  def qr_factorize(X, offset=0):
35      """
36      Compute QR factorization of X s.t. QR = X.

37

38      Parameters:
39          - X: square numpy ndarray.
40          - offset: (int) either 0 or 1. If offset is unity: compute
             Hessenberg-
41                      matrix.

42

43      Returns:
44          Q: square numpy ndarray, same shape as X. Rotation matrix.
45          R: square numpy ndarray, same shape as X. Upper triangular
             matrix if
46              offset is 0, Hessenberg-matrix if offset is 1.
47      """
48      assert offset in [0, 1]
49      assert type(X) == np.ndarray
50      assert X.shape[0] == X.shape[1]
```

```
51
52     R = copy.deepcopy(X)
53     Q = np.eye(X.shape[0])
54
55     for i in range(X.shape[0]-offset):
56         Pi = np.eye(R.shape[0])
57         _, Qi = hreflect1D(R[i+offset:, i])
58         Pi[i+offset:, i+offset:] = Qi
59
60         Q = Pi.dot(Q)
61         R = Pi.dot(R)
62
63     return Q.T, R
```

```
1  """
2  Algorithms for solving eigenvalue problems.
3
4  1. Compute diagonalization of 2x2 matrices via jacobi iteration.
5  2. Generalize Jacobi iteration for symmetric matrices.
6  """
7  import numpy as np
8  import copy
9  import warnings
10 from scipy import linalg as lin
11 from algorithms import helpers
12
13
14 def jacobi2x2(A):
15     """
16     Diagonalize a 2x2 matrix through jacobi step.
17
18     Solve: U' A U = E s.t. E is a diagonal matrix.
19
20     Parameters:
21         A - 2x2 numpy array.
22     Returns:
```

```
23            A - 2x2 diagonal numpy array
24        """
25        assert type(A) == np.ndarray
26        assert A.shape == (2, 2)
27        assert A[1, 0] == A[0, 1]
28
29        alpha = 0.5 * np.arctan(2*A[0, 1]/(A[1, 1] - A[0, 0]))
30        U = np.array([[np.cos(alpha), np.sin(alpha)],
31                       [-np.sin(alpha), np.cos(alpha)]])
32        E = np.matmul(U.T, np.matmul(A, U))
33        return E
34
35
36    def jacobi(X, precision=1e-6, debug=False):
37        """
38        Compute Eigenvalues and Eigenvectors for symmetric matrices.
39
40        Parameters:
41            X - 2D numpy ndarray which represents a symmetric matrix
42            precision - float in (0, 1). Convergence criterion.
43
44        Returns:
45            A - 1D numpy array with eigenvalues sorted by absolute value
46            U - 2D numpy array with associated eigenvectors (column).
47        """
48        assert 0 < precision < 1.
49        assert type(X) == np.ndarray
50        n, m = X.shape
51        assert n == m
52        assert all(np.isclose(X - X.T, np.zeros(n)).flatten())
53        A = copy.deepcopy(X)
54        U = np.eye(A.shape[0])
55        L = np.array([1])
56        iterations = 0
57
```

```python
58      while L.max() > precision:
59          L = np.abs(np.tril(A, k=0) - np.diag(A.diagonal()))
60          i, j = np.unravel_index(L.argmax(), L.shape)
61          alpha = 0.5 * np.arctan(2*A[i, j] / (A[i, i]-A[j, j]))
62
63          V = np.eye(A.shape[0])
64          V[i, i], V[j, j] = np.cos(alpha), np.cos(alpha)
65          V[i, j], V[j, i] = -np.sin(alpha), np.sin(alpha)
66
67          A = np.dot(V.T, A.dot(V))
68          U = U.dot(V)
69          iterations += 1
70
71      # Sort by eigenvalue (descending order) and flatten A
72      A = np.diag(A)
73      order = np.abs(A).argsort()[::-1]
74      if debug:
75          return iterations
76
77      return A[order], U[:, order]
78
79
80  def qrm(X, maxiter=15000, debug=False):
81      """
82      Compute Eigenvalues and Eigenvectors using the QR-Method.
83
84      Parameters:
85          - X: square numpy ndarray.
86      Returns:
87          - Eigenvalues of A.
88          - Eigenvectors of A.
89      """
90      n, m = X.shape
91      assert n == m
92
```

```python
93      # First stage: transform to upper Hessenberg-matrix.
94      A = copy.deepcopy(X)
95      conv = False
96      k = 0
97
98      # Second stage: perform QR-transformations.
99      while (not conv) and (k < maxiter):
100         k += 1
101         Q, R = helpers.qr_factorize(A)
102         A = R.dot(Q)
103
104         conv = np.alltrue(np.isclose(np.tril(A, k=-1), np.zeros((n, n)
                )))
105
106     if not conv:
107         warnings.warn("Convergence was not reached. Consider raising
                maxiter.")
108     if debug:
109         return k
110     Evals = A.diagonal()
111     order = np.abs(Evals).argsort()[::-1]
112     return Evals[order], Q[order, :]
113
114
115 def qrm2(X, maxiter=15000, debug=False):
116     """
117     First compute similar matrix in Hessenberg form, then compute the
118     Eigenvalues and Eigenvectors using the QR-Method.
119
120     Parameters:
121         - X: square numpy ndarray.
122     Returns:
123         - Eigenvalues of A.
124         - Eigenvectors of A.
125     """
```

```python
126      n, m = X.shape
127      assert n == m
128
129      # First stage: transform to upper Hessenberg-matrix.
130      A = lin.hessenberg(X)
131      conv = False
132      k = 0
133
134      # Second stage: perform QR-transformations.
135      while (not conv) and (k < maxiter):
136          k += 1
137          Q, R = helpers.qr_factorize(A)
138          A = R.dot(Q)
139
140          conv = np.alltrue(np.isclose(np.tril(A, k=-1), np.zeros((n, n)
             )))
141
142      if not conv:
143          warnings.warn("Convergence was not reached. Consider raising
              maxiter.")
144      if debug:
145          return k
146      Evals = A.diagonal()
147      order = np.abs(Evals).argsort()[::-1]
148      return Evals[order], Q[order, :]
149
150
151  def qrm3(X, maxiter=15000, debug=False):
152      """
153      First compute similar matrix in Hessenberg form, then compute the
154      Eigenvalues and Eigenvectors using the QR-Method.
155
156      Parameters:
157          - X: square numpy ndarray.
158      Returns:
```

```
159          - Eigenvalues of A.
160          - Eigenvectors of A.
161      """
162      n, m = X.shape
163      assert n == m
164
165      # First stage: transform to upper Hessenberg-matrix.
166      T = lin.hessenberg(X)
167
168      conv = False
169      k = 0
170
171      # Second stage: perform QR-transformations.
172      while (not conv) and (k < maxiter):
173          k += 1
174          Q, R = helpers.qr_factorize(T - T[n-1, n-1] * np.eye(n))
175          T = R.dot(Q) + T[n-1, n-1] * np.eye(n)
176
177          conv = np.alltrue(np.isclose(np.tril(T, k=-1), np.zeros((n, n)
             )))
178
179      if not conv:
180          warnings.warn("Convergence was not reached. Consider raising
             maxiter.")
181      if debug:
182          return k
183      Evals = T.diagonal()
184      order = np.abs(Evals).argsort()[::-1]
185      return Evals[order], Q[order, :]
```

6.4   Analysis: Figures

```python
import os

import copy

import pandas as pd

import numpy as np

import seaborn as sns

from scipy import linalg as lin

from scipy.stats import ortho_group

from matplotlib import pyplot as plt


datadir = os.path.join("analysis", "benchmarks.csv")

outpath = os.path.join("media", "plots")

trials = pd.read_csv(datadir, index_col=0)


trials.groupby(["algorithm", "dimension"]).iterations.describe()


# Boxplot iteration:
fig = plt.figure(figsize=(10, 5))

sns.boxplot(x="dimension", y="iterations", hue="algorithm", data=
    trials)

plt.yscale("log")

plt.title("Iterations needed before Convergence")

plt.savefig(os.path.join(outpath, "iterations_boxplot.png"))

plt.show()

plt.close()


# Boxplot elapsed time:
fig = plt.figure(figsize=(10, 5))

sns.boxplot(x="dimension", y="time", hue="algorithm", data=trials)

plt.title("Time needed before Convergence")

plt.ylabel("time (sec)")

plt.yscale('log')

plt.savefig(os.path.join(outpath, "time_boxplot.png"))

plt.show()

plt.close()
```

```
34
35  # Visualize Algorithm-Progress:
36  np.random.seed(42)
37  size = 5
38  Lambda = np.diag(np.random.randint(low=0, high=10, size=size))
39  G = ortho_group.rvs(dim=size)
40  X = np.dot(G, Lambda.dot(G.T))
41
42
43  def plot_factory(func):
44      def plotter(savepath, **fig_kw):
45          def algorithm_generator(*args, **kwargs):
46              return func(*args, **kwargs)
47
48          fig, ax = plt.subplots(nrows=2, ncols=2, **fig_kw)
49          algorithm_iterator = algorithm_generator()
50          j = -1
51
52          for i, A in enumerate(algorithm_iterator):
53              if i in (0, 1, 10, 75):
54                  j += 1
55
56                  hm = ax[j // 2, j % 2].imshow(A,
57                                                cmap=plt.get_cmap('
                                                    seismic'),
58                                                vmin=-X.max(),
59                                                vmax=X.max())
60                  ax[j // 2, j % 2].set_yticks([])
61                  ax[j // 2, j % 2].set_xticks([])
62                  ax[j // 2, j % 2].set_title("Iteration: " + str(i))
63
64                  if i > 75:
65                      break
66
67          fig.subplots_adjust(right=0.8)
```

29

```python
68          cbar_ax = fig.add_axes([0.85, 0.15, 0.05, 0.7])
69          fig.colorbar(hm, cax=cbar_ax)
70
71          sup_title = "Demonstrate {} on a {}x{} matrix".format(
72              func.__name__,
73              *X.shape)
74
75          fig.suptitle(sup_title)
76          fig.savefig(savepath)
77
78          return fig, ax
79
80      return plotter
81
82
83  @plot_factory
84  def jacobi():
85      """
86      Compute Eigenvalues and Eigenvectors for symmetric matrices using
         the
87      jacobi method.
88
89      Yields:
90          * A - 2D numpy array of current iteration step.
91      """
92      A = copy.deepcopy(X)
93      U = np.eye(A.shape[0])
94      L = np.array([1])
95      iterations = 0
96
97      while iterations < 5000:
98          L = np.abs(np.tril(A, k=0) - np.diag(A.diagonal()))
99          i, j = np.unravel_index(L.argmax(), L.shape)
100         alpha = 0.5 * np.arctan(2*A[i, j] / (A[i, i]-A[j, j]))
101
```

```python
102         V = np.eye(A.shape[0])
103         V[i, i], V[j, j] = np.cos(alpha), np.cos(alpha)
104         V[i, j], V[j, i] = -np.sin(alpha), np.sin(alpha)
105
106         A = np.dot(V.T, A.dot(V))
107         U = U.dot(V)
108         iterations += 1
109         yield A
110
111
112 @plot_factory
113 def qrm1():
114     """
115     Create generator for transformed matrices after applying the QR-
         Method.
116
117     Yields:
118         - T: 2D-numpy array. Similar matrix to X.
119     """
120     # First stage: transform to upper Hessenberg-matrix.
121     T = copy.deepcopy(X)
122
123     k = 0
124     # Second stage: perform QR-transformations.
125     while k < 5000:
126         k += 1
127         Q, R = np.linalg.qr(T)
128         T = R.dot(Q)
129         yield T
130
131
132 @plot_factory
133 def qrm2():
134     """
```

```python
135      Create generator for transformed matrices after applying the QR-
          Method.

136

137      Yields:

138          - T: 2D-numpy array. Similar matrix to X.

139      """

140      # First stage: transform to upper Hessenberg-matrix.

141      T = lin.hessenberg(X)

142

143      k = 0

144      # Second stage: perform QR-transformations.

145      while k < 5000:

146          if k == 0:

147              yield X

148          k += 1

149          Q, R = np.linalg.qr(T)

150          T = R.dot(Q)

151          yield T

152

153

154  @plot_factory

155  def qrm3():

156      """

157      First compute similar matrix in Hessenberg form, then compute the

158      Eigenvalues and Eigenvectors using the accelerated QR-Method.

159

160      Yields:

161          * T - 2D numpy array of current iteration step.

162      """

163      # First stage: transform to upper Hessenberg-matrix.

164      T = lin.hessenberg(X)

165      k = 0

166      n, _ = X.shape

167

168      # Second stage: perform QR-transformations.
```

```python
169     while k < 5000:
170         if k == 0:
171             yield X
172         k += 1
173         Q, R = np.linalg.qr(T - T[n-1, n-1] * np.eye(n))
174         T = R.dot(Q) + T[n-1, n-1] * np.eye(n)
175
176         yield T
177
178
179 jacobi(os.path.join(outpath, "jacobi.png"))
180 qrm1(os.path.join(outpath, "qrm1.png"))
181 qrm2(os.path.join(outpath, "qrm2.png"))
182 qrm3(os.path.join(outpath, "qrm3.png"))
183
184 plt.show()
185 plt.close()
```

6.5   Analysis: Unit tests

```python
"""
Automated tests for different algorithms.
"""
import os
import sys
import numpy as np
from threading import Thread
import pandas as pd
from algorithms import eigen
from scipy.stats import ortho_group
from tqdm import tqdm
from functools import wraps


data_out = os.path.join("data", "accuracy_tests.csv")


class AlgoTest(object):
    tests = {"algorithm": [],
             "dimension": [],
             "maxiter": [],
             "failed": []}

    def __init__(self, algo, dim, filepath, n_tests=1000, jobs=1,
                 *args, **kwargs):
        assert n_tests % jobs == 0
        self.algorithm = self.__get_test_algorithm(algo, *args, **
          kwargs)
        self.dim = dim
        self.n = n_tests // jobs
        self.failed = []
        self.jobs = jobs
        self.result = None
        self.path = filepath
```

```python
34          self.maxiter = kwargs.get("maxiter", None)

35

36          if not os.path.exists(self.path):
37              self.save(header=["algorithm", "dimension", "maxiter", "
                    failed"])

38

39      def __get_test_algorithm(self, algorithm, *args, **kwargs):
40          @wraps(algorithm)
41          def algo(X):
42              return algorithm(X, *args, **kwargs)
43          return algo

44

45      def __get_test_matrix(self):
46          """Return matrix with assosiated Eigenvalues."""
47          eigenvalues = np.random.uniform(size=self.dim)
48          eigenvectors = ortho_group.rvs(dim=self.dim)
49          Lambda = np.diag(eigenvalues)

50

51          matrix = np.dot(eigenvectors, Lambda).dot(eigenvectors.T)

52

53          order = np.abs(eigenvalues).argsort()[::-1]
54          return matrix, eigenvalues[order]

55

56      def __run_test(self):
57          """Run singular test."""
58          mat, true_eig = self.__get_test_matrix()
59          test_eig, _ = self.algorithm(mat)
60          test_res = np.alltrue(np.isclose(true_eig, test_eig))
61          self.failed.append(not test_res)

62

63      def __run_tests(self):
64          """Run multiple tests."""
65          for _ in range(self.n):
66              try:
67                  self.__run_test()
```

```python
68              except (KeyboardInterrupt, SystemExit):
69                  self.__save()
70                  sys.exit(0)

72      def run(self):
73          """Distribute tests accross threads."""

75          threadlist = [None] * self.jobs

77          for i in range(self.jobs):
78              threadlist[i] = Thread(target=self.__run_tests, daemon=
                  True)
79              threadlist[i].start()

81          for thread in threadlist:
82              thread.join()

84          self.result = sum(self.failed)
85          self.tests["algorithm"].append(self.algorithm.__name__)
86          self.tests["dimension"].append(self.dim)
87          self.tests["failed"].append(self.result)
88          if self.algorithm.__name__ == "jacobi":
89              self.tests["maxiter"].append(None)
90          else:
91              self.tests["maxiter"].append(self.maxiter)

93      def save(self, header=False):
94          if os.path.exists(self.path):
95                  print("Saving results.")
96          df = pd.DataFrame(self.tests)
97          with open(self.path, 'a') as f:
98              df.to_csv(f, header=header, index=False,
99                      columns=["algorithm", "dimension", "maxiter", "
                      failed"])
100
```

36

```
101
102  # Unit tests
103  if __name__ == "__main__":
104      # Define Flags.
105      JOBS = 20
106      MAXITER = (10, 100, 1000, 10000, 100000)
107      DIMS = range(3, 8)
108      ALGOS = {'jacobi': eigen.jacobi,
109               'qrm': eigen.qrm,
110               'qrm2': eigen.qrm2,
111               'qrm3': eigen.qrm3}
112
113      # Define parameters of all runs.
114      parameters = []
115      for algo in ALGOS.values():
116          for maxiter in MAXITER:
117              for dim in range(3, 8):
118                  param = (algo,
119                           maxiter if algo.__name__ != "jacobi" else
                                None,
120                           dim)
121                  parameters.append(param)
122
123      # Check progress of previous runs.
124      if os.path.exists(data_out):
125          print("Check progress:")
126          required = [(algo.__name__, m, dim) for (algo, m, dim) in
                parameters]
127          required = set(required)
128
129          progress = pd.read_csv(data_out)
130
131          done = []
132          for (a, d, m, _) in progress.values:
133              if np.isnan(m):
```

37

```python
134                    param = (a, None, d)
135                else:
136                    param = (a, int(m), d)
137                done.append(param)
138        done = set(done)
139
140        to_do = required.difference(done)
141        parameters = [(ALGOS[a], m, d) for a, m, d in to_do]
142
143    for algo, maxiter, dim in tqdm(parameters):
144        if algo is eigen.jacobi:
145            algo_test = AlgoTest(filepath=data_out,
146                                 algo=algo,
147                                 dim=dim,
148                                 jobs=JOBS)
149        else:
150            algo_test = AlgoTest(filepath=data_out,
151                                 algo=algo,
152                                 dim=dim,
153                                 maxiter=maxiter,
154                                 jobs=JOBS)
155        algo_test.run()
156
157    algo_test.save()
158
159    results = pd.read_csv(data_out)
160    results.head(10)
161    results.loc[results.algorithm == 'jacobi', 'maxiter'] = '-'
162
163    crosstab = pd.crosstab(index=[results.algorithm, results.maxiter],
164                           columns=results.dimension,
165                           values=results.failed,
166                           aggfunc=np.mean,
167                           dropna=True)
168    print(crosstab.to_latex())
```

7   References

[Börm and Mehl, 2012] Börm, S. and Mehl, C. (2012). *Numerical Methods for Eigenvalue Problems.* Walter de Gruyter GmbH & Co.KG, Berlin/Boston.

[Duda et al., 2000] Duda, R. O., Hart, P. E., and Stork, D. G. (2000). *Pattern Classification (2Nd Edition).* Wiley-Interscience, New York, NY, USA.

[Gentle, 1998] Gentle, J. E. (1998). *Numerical Linear Algebra for Applications in Statistics.* Springer Science and Business Media, New York.

[Härdle and Simar, 2015] Härdle, W. K. and Simar, L. (2015). *Applied Multivariate Statistical Analysis.* Springer-Verlag Gmbh, Berlin, Heidelberg.

[Hunter, 2007] Hunter, J. D. (2007). Matplotlib: A 2d graphics environment. *Computing In Science & Engineering*, 9(3):90–95.

[McKinney, 2010] McKinney, W. (2010). Data structures for statistical computing in python. In van der Walt, S. and Millman, J., editors, *Proceedings of the 9th Python in Science Conference*, pages 51 – 56.

[van Rossum, 1995] van Rossum, G. (1995). Python tutorial. Technical Report CS-R9526, Centrum voor Wiskunde en Informatica (CWI), Amsterdam.